

Polynomial Regression



GR 5205 / GU 4205
Section 3

Columbia University
Xiaofei Shi





So far...



We predict a scalar random variable Y as a linear function of $p-1$ different predictor variables X , plus noise: $\mathbf{Y} = \mathbf{X}\beta + \epsilon$

- Uncorrelated noise: unbiased estimator
- Gaussian noise: sampling distribution, hypothesis testing used in all packages
- $\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$

All results are based on the linear relationship holds.

What if the ground truth is something different?

Does the linear relationship holds?...



Solution: adding curvature!



Adding Curvature: Polynomial Regression

- If the relationship between Y and X is non-linear, we could try to capture that fact with a polynomial. For example:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_1^2 + \dots \beta_d X_1^d + \beta_{d+1} X_2 + \dots \beta_{p+d-1} X_p + \epsilon$$

- Instead of Y being linearly related to X1, it's polynomially related, with the degree of the polynomial being d
- Treat $x_1^2, x_1^3, \dots, x_1^d$ as additional “predictors” and include them in the design matrix X.
- Estimators are of the same form!

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$



Realization: polynomial degree = 2

$$a_0 + a_1 x + a_2 x^2$$

- In R:

```
out = lm(y ~ poly(x,2))
```

- In Python:

```
x = np.append(x, (x[:,1]**2).reshape(-1,1), 1)
x = statsmodels.tools.tools.add_constant(x)
model = sm.OLS(y, x).fit()
```

Or

```
model = np.poly1d(np.polyfit(x, y, 2))
```




Using a polynomial with degree = 2

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:          0.487
Model:                  OLS    Adj. R-squared:     0.482
Method:                 Least Squares    F-statistic:    93.69
Date:                  Sat, 17 Oct 2020    Prob (F-statistic): 2.54e-29
Time:                  18:15:47    Log-Likelihood:  -279.21
No. Observations:      200    AIC:          564.4
Df Residuals:          197    BIC:          574.3
Df Model:              2
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	1.3041	0.087	15.006	0.000	1.133	1.475
x1	0.0921	0.070	1.316	0.190	-0.046	0.230
x2	0.6920	0.052	13.397	0.000	0.590	0.794

```
=====
Omnibus:              0.242    Durbin-Watson:      1.895
Prob(Omnibus):        0.886    Jarque-Bera (JB):    0.184
Skew:                 -0.074    Prob(JB):            0.912
Kurtosis:             2.992    Cond. No.            2.45
=====
```






Using a polynomial with order = 3

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:          0.488
Model:                  OLS    Adj. R-squared:      0.480
Method:                 Least Squares    F-statistic:      62.31
Date:                   Sat, 17 Oct 2020    Prob (F-statistic): 2.46e-28
Time:                   18:18:04    Log-Likelihood:    -279.08
No. Observations:      200    AIC:              566.2
Df Residuals:          196    BIC:              579.4
Df Model:               3
Covariance Type:       nonrobust
=====
```


	coef	std err	t	P> t	[0.025	0.975]
const	1.3042	0.087	14.979	0.000	1.132	1.476
x1	0.0497	0.110	0.453	0.651	-0.167	0.266
x2	0.6923	0.052	13.376	0.000	0.590	0.794
x3	0.0150	0.030	0.503	0.616	-0.044	0.074

```
=====
Omnibus:                0.214    Durbin-Watson:          1.900
Prob(Omnibus):           0.898    Jarque-Bera (JB):        0.148
Skew:                    -0.067    Prob(JB):                0.929
Kurtosis:                2.999    Cond. No.                 6.07
=====
```




How to choose the polynomials?

- Smoothness:
Polynomials are very smooth, meaning that they and all their derivatives exist and are continuous.
Desirable if you are looking for a smooth dependence, not if there are sharp threshold or jumps.
Notice that one **can** approximate thresholds as accurate as one wants to, but ending up with very high order polynomials.



How to choose the polynomials?


- Overfitting:
A polynomial with degree d can fit any $d+1$ points. Using a high-order polynomial, or even summing a large number of low-order polynomials, can therefore lead to curves which come very close to the data we used to estimate them, but predict very badly.



Runge's Phenomenon

How to choose the polynomials?

- Picking the polynomial order:
 - scientific theory
 - carefully examining the diagnostics plots
 - variable and model selections





Other choices: Orthogonal Polynomials

Suppose that $x \in [-1, 1]$

- $f_0(x) = 1, \quad f_1(x) = x, \quad f_2(x) = x^2, \quad f_3(x) = x^3, \dots$

- Legendre polynomials:

$$g_0(x) = 1, \quad g_1(x) = x, \quad g_2(x) = \frac{1}{2}(3x^2 - 1), \quad g_3(x) = \frac{1}{2}(5x^3 - 3x), \dots$$

- gives the same results as the former simple polynomials;
- least squares optimization results are more stable and the standard errors of the coefficients are smaller



Other choices: beyond polynomials

- We are treating different powers of X as new features, and we can of course treat different functions of X as new features as well.
 - Fourier family: sines and cosines
 - ReLU and other activation functions
- Choose the functions:
 - scientific theory
 - carefully examining the diagnostics plots
 - variable and model selections



More to think about...

- Global trend v.s. local accuracy: a trade-off
- Piecewise polynomials, i.e. splines, are widely used in interpolation and fitting to avoid Runge's phenomenon, but more parameters need to be estimated.
- Outliers