



Modern Regression

STAT5241 Section 2

Statistical Machine Learning

Xiaofei Shi




Tasks for supervised learning

Input → Regressor → Predict real number




Input → Classifier → Predict category

Regression:




Regression:

- In everyday life we need to make decisions by taking into account lots of factors
- The question is what weight we put on each of these factors (how important are they with respect to the others)
- Suppose your task is to help build a evaluation system for food delivery apps

			
Available restaurants	30	10	20
Average delivery time	Next day	>3hr	1hr
Mandatory service fee	>10%	>20%	>13%
Score	9	7	8

Regression:

- Given an input x , we want to compute an output y
- For example:
 - predict Google's stock price using the current price of Bitcoin
 - predict arrival time using the traffic condition



Linear Regression:

- Given an input x , we want to compute an output y
- In linear regression we assume that y and x are related with the following


What we are trying to predict

Observed values

$$y = wx + \varepsilon$$

parameter we want to determine

noise term





Road map

- Build your model:

- 1) relationship: $\mathbf{y} = \mathbf{w}\mathbf{x} + \varepsilon$

- 2) preference: choose w to minimize $\arg \min_w \sum_i (y_i - wx_i)^2$

- Estimate your model parameters:

- 1) plugging in observed data to express your preference

- 2) get parameters estimation for your model


- Understand your model:

Linear Regression:

- Our goal is to estimate w from a training data of $\langle x_i, y_i \rangle$ pairs
- One easy way to determine w is to minimize the least squares error:

$$\arg \min_w \sum_i (y_i - wx_i)^2$$

- Why least squares?
 - easy to compute
 - has a nice probabilistic interpretation
- Several other approaches



If the noise is Gaussian with mean 0 then least squares is also the MLE of w

Solving linear regression using minimization

- Goal function:
- 3-step:
 - take derivative wrt parameter w
 - set it to 0
 - solve optimal w from the equation

$$\frac{\partial}{\partial w} \sum_i (y_i - wx_i)^2 = 2 \sum_i -x_i (y_i - wx_i) \Rightarrow$$


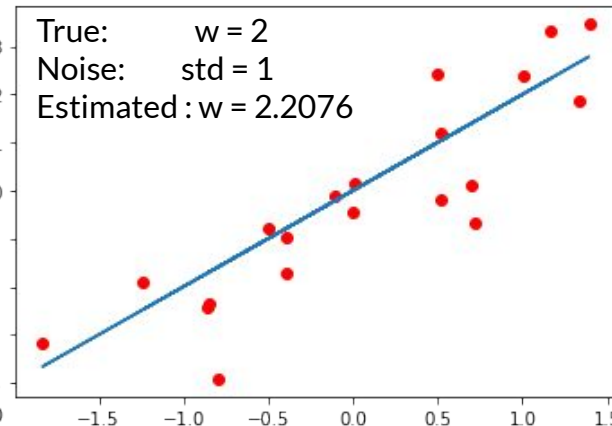

$$2 \sum_i x_i (y_i - wx_i) = 0 \Rightarrow$$

$$\sum_i x_i y_i = \sum_i wx_i^2 \Rightarrow$$

$$w = \frac{\sum_i x_i y_i}{\sum_i x_i^2}$$



Regression Example



Adding in intercept

- So far we assume the regression line passes through the origin
- What if the line does not?

$$y = w_0 + w_1x + \varepsilon$$

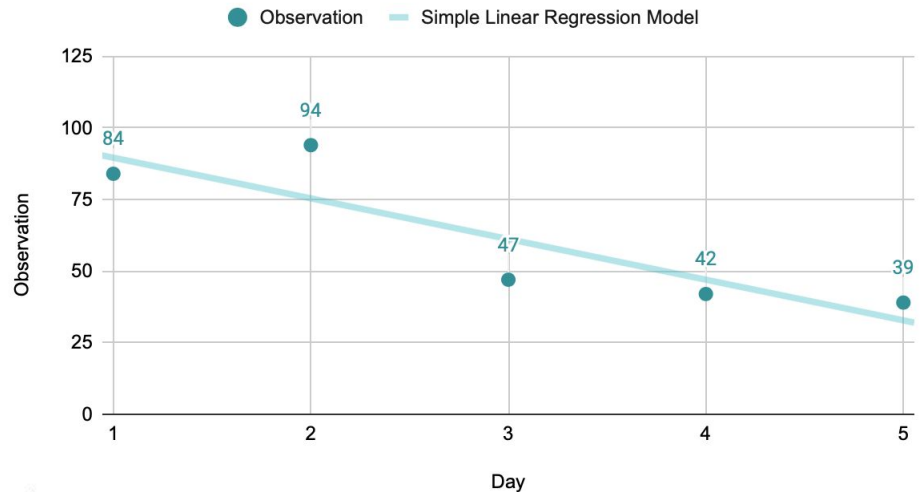
Adding in intercept!

- We can determine $w = (w_0, w_1)$ explicitly

$$w_0 = \frac{\sum_i y_i - w_1 x_i}{n}$$

$$w_1 = \frac{\sum_i x_i (y_i - w_0)}{\sum_i x_i^2}$$

Observation vs. Day



Multivariate Regression:

- What if we want to input more information?
- For example:
 - predict Google's stock price using the current price of Bitcoin might not be enough, we might want to also consider the stock price of Amazon, Apple, and other index
 - predict arrival time using the traffic condition, the weather condition, and the vehicle's condition to make it more accurate
- This becomes a multivariate linear regression problem:

$$y = w_0 + w_1x_1 + \dots + w_kx_k + \varepsilon$$

Google's price Bitcoin S&P 500



Multivariate Regression:

- What if we want to input more information?
- For example:
 - predict Google's stock price using the current price of Bitcoin might not be enough, we might want to also consider the stock price of Amazon, Apple, and other index
 - predict arrival time using the traffic condition, the weather condition, and the vehicle's condition to make it more accurate
- This becomes a multivariate linear regression problem:

$$y = w_0 + w_1x_1 + \dots + w_kx_k + \epsilon$$

Google's price Bitcoin S&P 500



Not all functions can be represented
using the input values directly!
How to capture nonlinearity?



How to capture nonlinearity?

- In some cases we would like to use polynomial or other terms based on the input data
 - Polynomial: $\phi_j(x) = x^j$ for $j=0 \dots n$
 - Gaussian: $\phi_j(x) = \frac{(x - \mu_j)}{2\sigma_j^2}$
 - Sigmoid: $\phi_j(x) = \frac{1}{1 + \exp(-s_j x)}$
- Are these still linear regression problems?



How to capture nonlinearity?

- In some cases we would like to use polynomial or other terms based on the input data
 - Polynomial: $\phi_j(x) = x^j$ for $j=0 \dots n$
 - Gaussian: $\phi_j(x) = \frac{(x - \mu_j)}{2\sigma_j^2}$
 - Sigmoid: $\phi_j(x) = \frac{1}{1 + \exp(-s_j x)}$
- Are these still linear regression problems?

As long as the coefficients are linear the equation is still a linear regression problem!

Nonlinear basis functions

- In some cases we would like to use polynomial or other terms based on the input data
 - Polynomial: $\phi_j(x) = x^j$ for $j=0 \dots n$
 - Gaussian: $\phi_j(x) = \frac{(x - \mu_j)}{2\sigma_j^2}$
 - Sigmoid: $\phi_j(x) = \frac{1}{1 + \exp(-s_j x)}$
- Linear regression can be applied in the same way to functions of these values
- As long as these functions can be directly computed from the observed values the parameters are still linear in the data and the problem remains a linear regression problem

Any function of the input values can be used. The solution for the parameters of the regression remains the same.



Nonlinear basis functions

- We use the new notation for the basis functions, linear regression can be written as

$$y = \sum_{j=0}^n w_j \phi_j(x)$$

- Nothing changed! Once again we can use 'least squares' to find the optimal solution to figure out parameter w

General linear regression problem

Our goal is to minimize the following loss function:

$$J(\mathbf{w}) = \sum_i (y^i - \sum_j w_j \phi_j(x^i))^2$$

$$y = \sum_{j=0}^K w_j \phi_j(x)$$

\mathbf{w} – vector of dimension $k+1$
 $\phi(x^i)$ – vector of dimension $k+1$
 y^i – a scalar

Moving to vector notations we get:

$$J(\mathbf{w}) = \sum_i (y^i - \mathbf{w}^T \phi(x^i))^2$$

We take the derivative w.r.t \mathbf{w}

$$\frac{\partial}{\partial \mathbf{w}} \sum_i (y^i - \mathbf{w}^T \phi(x^i))^2 = 2 \sum_i (y^i - \mathbf{w}^T \phi(x^i)) \phi(x^i)^T$$

Equating to 0 we get $2 \sum_i (y^i - \mathbf{w}^T \phi(x^i)) \phi(x^i)^T = 0 \Rightarrow$

$$\sum_i y^i \phi(x^i)^T = \mathbf{w}^T \left[\sum_i \phi(x^i) \phi(x^i)^T \right]$$

General linear regression problem

We take the derivative w.r.t \mathbf{w}

$$J(\mathbf{w}) = \sum_i (y^i - \mathbf{w}^T \phi(x^i))^2$$

$$\frac{\partial}{\partial \mathbf{w}} \sum_i (y^i - \mathbf{w}^T \phi(x^i))^2 = 2 \sum_i (y^i - \mathbf{w}^T \phi(x^i)) \phi(x^i)^T$$

Equating to 0 we get $2 \sum_i (y^i - \mathbf{w}^T \phi(x^i)) \phi(x^i)^T = 0 \Rightarrow$

$$\sum_i y^i \phi(x^i)^T = \mathbf{w}^T \left[\sum_i \phi(x^i) \phi(x^i)^T \right]$$

Define:


$$\Phi = \begin{pmatrix} \phi_0(x^1) & \phi_1(x^1) & \cdots & \phi_k(x^1) \\ \phi_0(x^2) & \phi_1(x^2) & \cdots & \phi_k(x^2) \\ \vdots & \vdots & \cdots & \vdots \\ \phi_0(x^n) & \phi_1(x^n) & \cdots & \phi_k(x^n) \end{pmatrix}$$

Then deriving w
we get:


$$\mathbf{w} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$




Example: polynomial regression




degree = 1, CV = 0.6



degree = 3, CV = 1.5



degree = 5, CV = 6.0



degree = 7, CV = 15.6

Thoughts ?

Recap of linear regression:

- Build your model:

1) relationship:
$$y = \sum_{j=0}^k w_j \phi_j(x)$$

2) preference: choose w to minimize
$$J(w) = \sum_i (y^i - \sum_j w_j \phi_j(x^i))^2$$

- Estimate your model parameters:

1) plugging in observed data to express your preference

2) get parameters estimation for your model
$$w = (\Phi^T \Phi)^{-1} \Phi^T y$$

- Understand your model

Potential problems:

- collinearity
- too many non-zero but very small coefficients
- too slow





Regularizer: ridge regression

- If $\Phi^T \Phi$ is not invertible, or its determinant is very small, the optimal w is not going to be stable
- n equations $<$ p unknowns – underdetermined system of linear equations many feasible solutions

Need to impose extra constraints!



Regularizer: ridge regression

- If $\Phi^T \Phi$ is not invertible, or its determinant is very small, the optimal w is not going to be stable
- n equations $<$ k unknowns – underdetermined system of linear equations many feasible solutions
- Adding in penalty term into loss function:

$$\begin{aligned} J(\beta) &= \sum_i \left(y^i - \sum_j \beta_j \phi_j(x^i) \right)^2 + \lambda \sum_j \beta_j^2 \\ &= \|y - \Phi(x)\beta\|_2^2 + \lambda \|\beta\|_2^2 \end{aligned}$$

different norms of
matrix and vectors

- Equivalent to a MAP optimization problem \longrightarrow HW1

$$\hat{\beta} = (\Phi^T(x)\Phi(x) + \lambda I)^{-1} \Phi^T(x)y$$



Regularizer: ridge regression

- If $\Phi^T \Phi$ is not invertible, or its determinant is very small, the optimal w is not going to be stable
- n equations < k unknowns – underdetermined system of linear equations many feasible solutions
- Adding in penalty term into loss function:

$$\begin{aligned} J(\beta) &= \sum_i \left(y^i - \sum_j \beta_j \phi_j(x^i) \right)^2 + \lambda \sum_j \beta_j^2 \\ &= \|y - \Phi(x)\beta\|_2^2 + \lambda \|\beta\|_2^2 \end{aligned}$$

different norms of
matrix and vectors

- Equivalent to a MAP optimization problem \longrightarrow HW1

$$\hat{\beta} = (\Phi^T(x)\Phi(x) + \lambda I)^{-1} \Phi^T(x)y$$


- Don't have to worry about invertibility anymore!



Regularizer: ridge regression

Ridge Regression:

$$\text{pen}(\beta) = \|\beta\|_2^2$$





Regularizer: lasso

- n equations $<$ k unknowns – underdetermined system of linear equations many feasible solutions
- Sometimes our goal is to learn a ***sparse*** representation: select the most useful features!
- How to achieve?



Regularizer: lasso

- n equations $<$ k unknowns – underdetermined system of linear equations many feasible solutions
- Sometimes our goal is to learn a ***sparse*** representation: select the most useful features!
- How to achieve?

$$J(\beta) = \|y - \Phi(x)\beta\|_2^2 + \lambda\|\beta\|_0$$



No closed form!
Hard to solve!



Regularizer: lasso

- n equations $<$ k unknowns – underdetermined system of linear equations many feasible solutions
- Sometimes our goal is to learn a ***sparse*** representation: select the most useful features!
- How to achieve?


$$J(\beta) = \|y - \Phi(x)\beta\|_2^2 + \lambda\|\beta\|_1$$




No closed form!
Getting easier!

Lasso or Ridge?

Lasso Coefficients



Ridge Coefficients




Ridge Regression:

$$\text{pen}(\beta) = \|\beta\|_2^2$$

Lasso:

$$\text{pen}(\beta) = \|\beta\|_1$$

Ideally l0 penalty,
but optimization
becomes non-convex



Lasso (l_1 penalty) results in sparse solutions – vector with more zero coordinates
Good for high-dimensional problems – don't have to store all coordinates, interpretable solution!

Regression to classification

- Instead of giving scores to these apps, can you tell which app to use?
- Can we predict the “probability” of class label – a real number – using regression methods?
- But output (probability) needs to be in $[0,1]$

A way to make categorical variables continuous!

			
Available restaurants	30	10	20
Average delivery time	Next day	>3hr	1hr
Mandatory service fee	>10%	>20%	>13%
Score	9	7	8

Logistic regression


- Instead of modeling $Y = 0$ or 1 directly, we modify the probability of $P(Y=0|x)$ as

$$P(Y = 0|X) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

Logistic function applied to a linear function of the data

Logistic
function
(or Sigmoid): $\frac{1}{1 + \exp(-z)}$

Features can be discrete or continuous!



2 categories

Assumes the following functional form for $P(Y|X)$:

$$P(Y = 0|X) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$\Rightarrow P(Y = 1|X) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$\Rightarrow \frac{P(Y = 1|X)}{P(Y = 0|X)} = \exp(w_0 + \sum_i w_i X_i) \geq 1$$

$$\Rightarrow w_0 + \sum_i w_i X_i \geq 0$$

2 categories

Assumes the following functional form for $P(Y|X)$:


$$P(Y = 0|X) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

Decision boundary: Note - Labels are 0,1

$$P(Y = 0|X) \geq P(Y = 1|X)$$

$$w_0 + \sum_i w_i X_i \geq 0$$

(Linear Decision Boundary)





Expressing conditional likelihood

$$P(Y = 0|\mathbf{X}, \mathbf{w}) = \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)}$$



$$P(Y = 1|\mathbf{X}, \mathbf{w}) = \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)}$$

$$l(\mathbf{w}) = \sum_j \left[y^j (w_0 + \sum_i w_i x_i^j) - \ln(1 + \exp(w_0 + \sum_i w_i x_i^j)) \right]$$






Expressing conditional likelihood

$$\begin{aligned} P(Y = 0|\mathbf{X}, \mathbf{w}) &= \frac{1}{1 + \exp(w_0 + \sum_i w_i X_i)} \\ P(Y = 1|\mathbf{X}, \mathbf{w}) &= \frac{\exp(w_0 + \sum_i w_i X_i)}{1 + \exp(w_0 + \sum_i w_i X_i)} \end{aligned} \quad \Longrightarrow \quad l(\mathbf{w}) = \sum_j \left[y^j (w_0 + \sum_i w_i x_i^j) - \ln(1 + \exp(w_0 + \sum_i w_i x_i^j)) \right]$$

- Bad: we cannot find explicit solution anymore
- Good: it is guaranteed to have a unique solution, and we can still solve this problem numerically

Convex optimization



A function $l(w)$ is called **concave** if the line joining two points $l(w_1), l(w_2)$ on the function does not go above the function on the interval $[w_1, w_2]$


(Strictly) Concave functions have a unique maximum!



Convex



Both Concave & Convex



Neither

Convex optimization for logistic regression

Gradient ascent rule for w_0 :

$$w_0^{(t+1)} \leftarrow w_0^{(t)} + \eta \left. \frac{\partial l(\mathbf{w})}{\partial w_0} \right|_t$$

$$l(\mathbf{w}) = \sum_j \left[y^j (w_0 + \sum_i^d w_i x_i^j) - \ln(1 + \exp(w_0 + \sum_i^d w_i x_i^j)) \right]$$

$$\frac{\partial l(\mathbf{w})}{\partial w_0} = \sum_j \left[y^j - \underbrace{\frac{1}{1 + \exp(w_0 + \sum_i^d w_i x_i^j)} \cdot \exp(w_0 + \sum_i^d w_i x_i^j)}_{\text{probability}} \right]$$

$$w_0^{(t+1)} \leftarrow w_0^{(t)} + \eta \sum_j [y^j - \hat{P}(Y^j = 1 \mid \mathbf{x}^j, \mathbf{w}^{(t)})]$$

Convex optimization for logistic regression

Gradient ascent algorithm: iterate until change $< \epsilon$

$$w_0^{(t+1)} \leftarrow w_0^{(t)} + \eta \sum_j [y^j - \hat{P}(Y^j = 1 \mid \mathbf{x}^j, \mathbf{w}^{(t)})]$$

For $i=1, \dots, d$,

$$w_i^{(t+1)} \leftarrow w_i^{(t)} + \eta \sum_j x_i^j [y^j - \underbrace{\hat{P}(Y^j = 1 \mid \mathbf{x}^j, \mathbf{w}^{(t)})}_{\text{Predict what current weight thinks label Y should be}}]$$

repeat

- Gradient ascent is simplest of optimization approaches
 - e.g., Newton method, Conjugate gradient ascent, IRLS (see Bishop 4.3.3)





More than 2 categories

- Logistic regression in more general case, where $Y \in \{y_1, \dots, y_K\}$

for $k < K$

$$P(Y = y_k | X) = \frac{\exp(w_{k0} + \sum_{i=1}^d w_{ki} X_i)}{1 + \sum_{j=1}^{K-1} \exp(w_{j0} + \sum_{i=1}^d w_{ji} X_i)}$$

for $k=K$ (normalization, so no weights for this class)

$$P(Y = y_K | X) = \frac{1}{1 + \sum_{j=1}^{K-1} \exp(w_{j0} + \sum_{i=1}^d w_{ji} X_i)}$$

Predict $f^*(x) = \arg \max_{Y=y} P(Y = y | X = x)$



More than 2 categories

- Logistic regression in more general case, where $Y \in \{y_1, \dots, y_K\}$

for $k < K$

$$P(Y = y_k | X) = \frac{\exp(w_{k0} + \sum_{i=1}^d w_{ki} X_i)}{1 + \sum_{j=1}^{K-1} \exp(w_{j0} + \sum_{i=1}^d w_{ji} X_i)}$$

for $k=K$ (normalization, so no weights for this class)

Are decision boundaries still linear? Why?

$$P(Y = y_K | X) = \frac{1}{1 + \sum_{j=1}^{K-1} \exp(w_{j0} + \sum_{i=1}^d w_{ji} X_i)}$$

Predict $f^*(x) = \arg \max_{Y=y} P(Y = y | X = x)$





References

- Christopher Bishop: Pattern Recognition and Machine Learning, Chapter 3, 4
- Kutner, Nachtsheim and Neter: Applied Linear Regression Models.
- Agresti: Foundations of Linear and Generalized Linear Models.
- Ziv Bar-Joseph, Tom Mitchell, Pradeep Ravikumar and Aarti Singh: CMU 10-701