

A short horizontal bar with a teal segment on the left and an orange segment on the right.

# Introduction to Artificial Neural Networks

STAT5241 Section 2

Statistical Machine Learning

Xiaofei Shi

## Images & Video



Product  
Recommendation


amazon




## Text & Language




Relational Data/  
Social Network




## Speech & Audio




## Gene Expression



fMRI



Tumor region



# Overview



- A neural network is a supervised learning method. It can be applied to both regression and classification problems.
- The main idea is to extract linear combinations of the inputs as derived features, and then model the target as a nonlinear function of these features.
- The nonlinear transformation contributes to the model flexibility.
- Today, we will focus on the most widely used "vanilla" neural net, also called the single hidden layer feedforward neural networks.




# Recall Logistic Regression

$$P(Y = 1|X) = \frac{1}{1 + \exp(-(w_0 + \sum_i w_i X_i))}$$

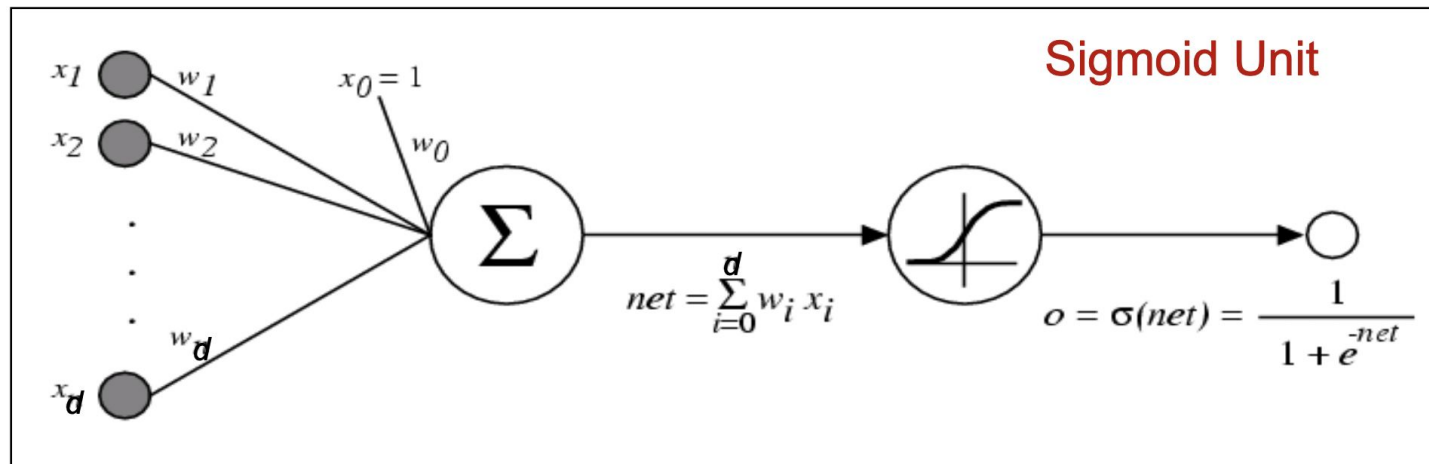
Logistic function applied to a linear function of the data

**Logistic function**  
**(or Sigmoid):**  $\frac{1}{1 + \exp(-z)}$




# Logistic function as a graph

$$\text{Output, } o(\mathbf{x}) = \sigma(w_0 + \sum_i w_i X_i) = \frac{1}{1 + \exp(-(w_0 + \sum_i w_i X_i))}$$




# Use neural networks to learn $f: X \rightarrow Y$


- $f$  can be a **non-linear** function
- $X$  (vector of) continuous and/or discrete variables
- $Y$  (**vector** of) continuous and/or discrete variables
- Neural networks - Represent  $f$  by network of logistic/sigmoid units:



# Use neural networks to learn $f: X \rightarrow Y$



Two layers of logistic units



Highly non-linear decision surface

# Consider humans:

- Neuron switching time  $\sim .001$  second
  - Number of neurons  $\sim 10^{10}$
  - Connections per neuron  $\sim 10^{4-5}$
  - Scene recognition time  $\sim .1$  second
  - 100 inference steps doesn't seem like enough
- much parallel computation

Properties of artificial neural nets (ANN's):

- Many neuron-like threshold switching units
- Many weighted interconnections among units
- Highly parallel, distributed process



# Overview

- ▶ Derived features  $Z_m$  are obtained by applying the *activation function*  $\sigma$  to linear combinations of the inputs:


$$Z_m = \sigma(\alpha_{0m} + \alpha_m^T X), m = 1, \dots, M.$$

- ▶ The target  $Y_k$  (or  $T_k$  in the figure) is modeled as a function of linear combinations of the  $Z_m$ :

$$T_k = \beta_{0k} + \beta_k^T Z, \quad k = 1, \dots, K.$$

- ▶ The output function  $g_k(T)$  allows a final transformation of the vector of outputs  $T$ :

$$f_k(X) = g_k(T), \quad k = 1, \dots, K.$$




Schematic of a single hidden layer,  
feed-forward neural network



# Artificial neurons

- Each artificial neuron has inputs and produces a single output which can be sent to multiple other neurons. The inputs can be the feature values of a sample of external data, such as images or documents, or they can be the outputs of other neurons.
- The outputs of the final output neurons of the neural net accomplish the task.



Schematic of a single hidden layer,  
feed-forward neural network

# Artificial neurons

- Neuron pre-activation (or input activation):

$$a(\mathbf{x}) = b + \sum_i w_i x_i = b + \mathbf{w}^\top \mathbf{x}$$

- Neuron output activation:


$$h(\mathbf{x}) = g(a(\mathbf{x})) = g(b + \sum_i w_i x_i)$$

where

$\mathbf{w}$  are the weights (parameters)


$b$  is the bias term

$g(\cdot)$  is called the activation function



# Activation functions

- An activation function of a node defines the output of that node given an input or set of inputs.
- Usually nonlinear.




Schematic of a single hidden layer,  
feed-forward neural network

# Activation functions


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$


## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$




# Activation function: linear

- No nonlinear transformation
- No input squashing




# Activation function: sigmoid

- Always positive.
- Squashing the neuron's output between 0 and 1.
- Strictly increasing.



# Activation function: tanh

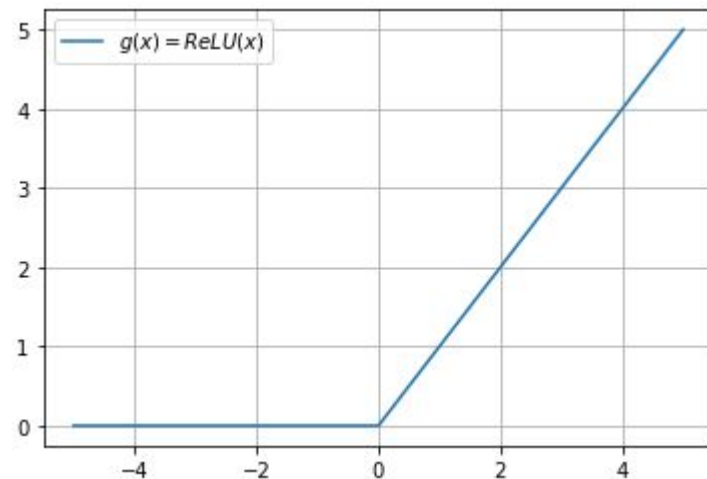
- Can be positive and negative.
- Squashing the neuron's output between -1 and 1.
- Strictly increasing.





# Activation function: sigmoid

- Always positive.
- Pushing the neuron's output above 0.
- Strictly increasing.



# Prediction using neural networks

**Prediction** – Given neural network (hidden units and weights), use it to predict the label of a test point

**Forward Propagation** –

Start from input layer

For each subsequent layer, compute output of sigmoid unit

Sigmoid unit:

$$o(\mathbf{x}) = \sigma(w_0 + \sum_i w_i x_i)$$

1-Hidden layer,  
1 output NN:

$$o(\mathbf{x}) = \sigma \left( w_0 + \sum_h w_h \underbrace{\sigma \left( w_0^h + \sum_i w_i^h x_i \right)}_{O_h} \right)$$

# Fitting neural networks: regression tasks

Recall our model is:

$$\begin{aligned}Z_m &= \sigma(\alpha_{0m} + \alpha_m^T X), \quad m = 1, \dots, M. \\T_k &= \beta_{0k} + \beta_k^T Z, \quad k = 1, \dots, K. \\f_k(X) &= g_k(T), \quad k = 1, \dots, K.\end{aligned}$$

The unknown parameters of the model are often called *weights*. We denote the complete set of weights by  $\theta$ , which consists of

$$\begin{aligned}\{\alpha_{0m}, \alpha_m; \quad m = 1, 2, \dots, M\} & \quad M(p+1) \text{ weights,} \\ \{\beta_{0k}, \beta_k; \quad k = 1, 2, \dots, K\} & \quad K(M+1) \text{ weights.}\end{aligned}$$

For regression, we use the squared error loss

$$R(\theta) = \sum_{k=1}^K \sum_{i=1}^n (y_{ik} - f_k(x_i))^2.$$

# Fitting neural networks: classification tasks

Recall our model is:

$$\begin{aligned}Z_m &= \sigma(\alpha_{0m} + \alpha_m^T X), \quad m = 1, \dots, M. \\T_k &= \beta_{0k} + \beta_k^T Z, \quad k = 1, \dots, K. \\f_k(X) &= g_k(T), \quad k = 1, \dots, K.\end{aligned}$$

The unknown parameters of the model are often called *weights*. We denote the complete set of weights by  $\theta$ , which consists of

$$\begin{aligned}\{\alpha_{0m}, \alpha_m; \quad m = 1, 2, \dots, M\} & \quad M(p+1) \text{ weights,} \\ \{\beta_{0k}, \beta_k; \quad k = 1, 2, \dots, K\} & \quad K(M+1) \text{ weights.}\end{aligned}$$

For classification we use either squared error or cross-entropy

$$R(\theta) = - \sum_{i=1}^n \sum_{k=1}^K y_{ik} \log f_k(x_i),$$

and the corresponding classifier is  $G(x) = \operatorname{argmax}_k f_k(x)$ .

# Connection to gradient descent

Assume we use squared error loss. Let  $z_{mi} = \sigma(\alpha_{0m} + \alpha_m^T x_i)$  and let  $z_i = (z_{1i}, z_{2i}, \dots, z_{Mi})$ . Then we have

$$R(\theta) \equiv \sum_{i=1}^n R_i = \sum_{i=1}^n \sum_{k=1}^K (y_{ik} - f_k(x_i))^2,$$

where

$$f_k(x_i) = g_k(\beta_{0k} + \beta_k^T z_i) = g_k \left( \beta_{0k} + \sum_{m=1}^M \beta_{km} \sigma(\alpha_{0m} + \alpha_m^T x_i) \right).$$

The derivatives are

$$\begin{aligned} \frac{\partial R_i}{\partial \beta_{km}} &= -2(y_{ik} - f_k(x_i)) g'_k(\beta_k^T z_i) z_{mi}, \\ \frac{\partial R_i}{\partial \alpha_{ml}} &= -\sum_{k=1}^K 2(y_{ik} - f_k(x_i)) g'_k(\beta_k^T z_i) \beta_{km} \sigma'(\alpha_m^T x_i) x_{il}, \end{aligned}$$



# Connection to gradient descent

Assume we use squared error loss. Let  $z_{mi} = \sigma(\alpha_{0m} + \alpha_m^T x_i)$  and let  $z_i = (z_{1i}, z_{2i}, \dots, z_{Mi})$ . Then we have

$$R(\theta) \equiv \sum_{i=1}^n R_i = \sum_{i=1}^n \sum_{k=1}^K (y_{ik} - f_k(x_i))^2,$$

where

$$f_k(x_i) = g_k(\beta_{0k} + \beta_k^T z_i) = g_k\left(\beta_{0k} + \sum_{m=1}^M \beta_{km} \sigma(\alpha_{0m} + \alpha_m^T x_i)\right).$$

The derivatives are

$$\begin{aligned} \frac{\partial R_i}{\partial \beta_{km}} &= -2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)z_{mi}, \\ \frac{\partial R_i}{\partial \alpha_{ml}} &= -\sum_{k=1}^K 2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)\beta_{km}\sigma'(\alpha_m^T x_i)x_{il}, \end{aligned}$$

# Updating rule

Assume we use squared error loss. Let  $z_{mi} = \sigma(\alpha_{0m} + \alpha_m^T x_i)$  and let  $z_i = (z_{1i}, z_{2i}, \dots, z_{Mi})$ . Then we have

$$R(\theta) \equiv \sum_{i=1}^n R_i = \sum_{i=1}^n \sum_{k=1}^K (y_{ik} - f_k(x_i))^2,$$

where

$$f_k(x_i) = g_k(\beta_{0k} + \beta_k^T z_i) = g_k\left(\beta_{0k} + \sum_{m=1}^M \beta_{km} \sigma(\alpha_{0m} + \alpha_m^T x_i)\right).$$

A gradient update at the  $(r+1)$ st iteration has the form

$$\begin{aligned}\beta_{km}^{(r+1)} &= \beta_{km}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial R_i}{\partial \beta_{km}^{(r)}}, \\ \alpha_{ml}^{(r+1)} &= \alpha_{ml}^{(r)} - \gamma_r \sum_{i=1}^N \frac{\partial R_i}{\partial \alpha_{ml}^{(r)}}.\end{aligned}$$

# Updating rule

If we write the gradients as

$$\frac{\partial R_i}{\partial \beta_{km}} = -2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)z_{mi},$$

$$\frac{\partial R_i}{\partial \alpha_{ml}} = -\sum_{k=1}^K 2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)\beta_{km}\sigma'(\alpha_m^T x_i)x_{il}.$$



# Back propagation

If we write the gradients as

$$\frac{\partial R_i}{\partial \beta_{km}} = -2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)z_{mi},$$

$$\frac{\partial R_i}{\partial \alpha_{ml}} = -\sum_{k=1}^K 2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)\beta_{km}\sigma'(\alpha_m^T x_i)x_{il}.$$

$$\begin{array}{l} \downarrow \quad \downarrow \\ \frac{\partial R_i}{\partial \beta_{km}} = \delta_{ki} z_{mi}, \\ \frac{\partial R_i}{\partial \alpha_{ml}} = s_{mi} x_{il}. \end{array}$$



# Back propagation

If we write the gradients as

$$\begin{aligned}\frac{\partial R_i}{\partial \beta_{km}} &= \delta_{ki} z_{mi}, \\ \frac{\partial R_i}{\partial \alpha_{ml}} &= s_{mi} x_{il}.\end{aligned}$$

In some sense,  $\delta_{ki}$  and  $s_{mi}$  are “errors” at the output and hidden layer units. The errors satisfy

$$s_{mi} = \sigma'(\alpha_m^T x_i) \sum_{k=1}^K \beta_{km} \delta_{ki}.$$

They are called the *back-propagation equations*. The updates can be implemented with a two-pass algorithm:

- ▶ *forward pass*: fix weights, compute the predicted values  $\hat{f}_k(x_i)$ .
- ▶ *backward pass*: errors  $\delta_{ki}$  are computed, and back-propagated to give the errors  $s_{mi}$ . Then use both sets of errors to compute the gradients.


# Back propagation

- Gradient descent over entire *network* weight vector
- Easily generalized to arbitrary directed graphs
- Will find a local, not necessarily global error minimum
  - In practice, often works well (can run multiple times)

- Often include weight *momentum*  $\alpha$

$$\Delta w_{i,j}(n) = \eta \delta_j x_{i,j} + \alpha \Delta w_{i,j}(n-1)$$

- Minimizes error over *training* examples
  - Will it generalize well to subsequent examples?
- Training can take thousands of iterations → slow!
- Using network after training is very fast




Schematic of a single hidden layer, feed-forward neural network



# Starting values

- ▶ If the weights are near zero, then the operative part of the sigmoid is roughly zero.
- ▶ Usually starting values for weights are chosen to be random values near zero.
- ▶ Hence the model starts out nearly linear, and becomes nonlinear as the weights increases.



**FIGURE 11.3.** Plot of the sigmoid function  $\sigma(v) = 1/(1 + \exp(-v))$  (red curve), commonly used in the hidden layer of a neural network. Included are  $\sigma(sv)$  for  $s = \frac{1}{2}$  (blue curve) and  $s = 10$  (purple curve). The scale parameter  $s$  controls the activation rate, and we can see that large  $s$  amounts to a hard activation at  $v = 0$ . Note that  $\sigma(s(v - v_0))$  shifts the activation threshold from 0 to  $v_0$ .

# Multiple minima



The error function  $R(\theta)$  is nonconvex, possessing many local minima.

The solution we obtained from back-propagation is a local minimum.

Usually, we try a number of random starting configuration, and choose the solution giving lowest error, or use the average predictions over the collection of networks as the final prediction.



# Multiple minima

- ▶ Often neural networks have too many weights and will overfit the data at the global minimum of  $R$ .
- ▶ A regularization method is *weight decay*. We add a penalty to the error function  $R(\theta) + \lambda J(\theta)$ , where


$$J(\theta) = \sum_{k,m} \beta_{km}^2 + \sum_{m,l} \alpha_{ml}^2.$$

- ▶  $\lambda \geq 0$  is a tuning parameter, can be chosen by cross-validation.


Neural Network - 10 Units, No Weight Decay



Neural Network - 10 Units, Weight Decay=0.02



# Summary: pipeline for simple layer neural net



# Expressive capability of ANNs

---

Boolean functions:

- Every boolean function can be represented by network with single hidden layer
- but might require exponential (in number of inputs) hidden units

Continuous functions:

- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer [Cybenko 1989; Hornik et al. 1989]
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers [Cybenko 1988].






# Example

target function:


Input	Output
10000000 →	10000000
01000000 →	01000000
00100000 →	00100000
00010000 →	00010000
00001000 →	00001000
00000100 →	00000100
00000010 →	00000010
00000001 →	00000001




# Example


Learned hidden layer representation:

Input		Hidden Values		Output
10000000	→	.89 .04 .08	→	10000000
01000000	→	.01 .11 .88	→	01000000
00100000	→	.01 .97 .27	→	00100000
00010000	→	.99 .97 .71	→	00010000
00001000	→	.03 .05 .02	→	00001000
00000100	→	.22 .99 .99	→	00000100
00000010	→	.80 .01 .98	→	00000010
00000001	→	.60 .94 .01	→	00000001



# Training





Typical input images

90% accurate learning head pose, and recognizing 1-of-20 faces



# Why training is hard




- Underfitting: use better optimization:
  - use better optimization tools (e.g. batch-normalization, 2nd-order methods).
  - use GPUs, distributed computing.
- Overfitting: use better regularization:
  - unsupervised pre-training
  - stochastic drop-out training
- For many large-scale practical problems, have to scale up:
  - ReLu nonlinearity
  - initialization (e.g. Kaiming He's initialization)
  - stochastic gradient descent
  - momentum, batch-normalization, and drop-out

# Preprocessing

- One-hot representation: class 0 or class 1  $\rightarrow (1,0)$  or  $(0,1)$
- Normalizing the inputs will speed up training (Lecun et al. 1998)
  - could normalization be useful at the level of the hidden layers?
- Batch normalization is an attempt to do that
  - each unit's pre-activation is normalized (mean subtraction, stddev division)
  - during training, mean and stddev is computed for each minibatch
  - backpropagation takes into account the normalization
  - at test time, the global mean / stddev is used

$$\mathbf{a}^{(k)}(\mathbf{x}) = \mathbf{b}^{(k)} + \mathbf{W}^{(k)}\mathbf{h}^{(k-1)}(\mathbf{x})$$



# Initialization of parameters

- Initialize biases to 0
- For weights
  - Can not initialize weights to 0 with tanh activation
    - All gradients would be zero (saddle point)
  - Can not initialize all weights to the same value
    - All hidden units in a layer will always behave the same
    - Need to break symmetry
  - Sample  $\mathbf{W}_{i,j}^{(k)}$  from  $U[-b, b]$ , where

$$b = \frac{\sqrt{6}}{\sqrt{H_k + H_{k-1}}}$$

Sample around 0 and  
break symmetry




Size of  $\mathbf{h}^{(k)}(\mathbf{x})$






# Overfitting

- Overfitting often occurs in applications of neural networks.
- Ways to overcome:
  - Early stopping:  
Stop training process early.
  - Dropout:  
Use random binary masks.




# Early stopping



# Dropouts

- Cripple neural network by removing hidden units stochastically
  - each hidden unit is set to 0 with probability 0.5
  - hidden units cannot co-adapt to other units
  - hidden units must be more generally useful
- Could use a different dropout probability, but 0.5 usually works well



# Model selection

- Training Protocol:
  - Train your model on the **Training Set**  $\mathcal{D}^{\text{train}}$
  - For model selection, use **Validation Set**  $\mathcal{D}^{\text{valid}}$ 
    - Hyper-parameter search: hidden layer size, learning rate, number of iterations/epochs, etc.
  - Estimate generalization performance using the **Test Set**  $\mathcal{D}^{\text{test}}$
- Remember: Generalization is the behavior of the model on **unseen examples**.

# Optimization

---

- SGD with momentum, batch-normalization, and dropout usually works very well
- Pick learning rate by running on a subset of the data
  - Start with large learning rate & divide by 2 until loss does not diverge
  - Decay learning rate by a factor of  $\sim 100$  or more by the end of training
  - Use ReLU nonlinearity
  - Initialize parameters so that each feature across layers has similar variance. Avoid units in saturation.
- Use adapted learning rate



# Summary:

---

- Actively used to model distributed computation in brain
- Highly non-linear regression/classification
- Vector-valued inputs and outputs
- Potentially millions of parameters to estimate - overfitting
- Hidden layers learn intermediate representations – how many to use?
- Prediction – Forward propagation
- Gradient descent (Back-propagation), local minima problems
- Coming back in new form as deep networks
  - Try different/deeper architecture



# References

- Christopher Bishop: Pattern Recognition and Machine Learning, Chapter 5
- Ziv Bar-Joseph, Tom Mitchell, Pradeep Ravikumar and Aarti Singh: CMU 10-701
- Ryan Tibshirani: CMU 10-725
- Ruslan Salakhutdinov: CMU 10-703
- <https://towardsdatascience.com/neural-network-architectures-156e5bad51ba>
- <https://towardsdatascience.com/applied-deep-learning-part-4-convolutional-neural-networks-584bc134c1e2>